

# Solution

December 10, 2017

## 0.1 Ibrahim AbuAlhaol, Ottawa, Canada

- DEC 10, 2017
- Q1: Using data in Bar.csv, How much faster does the treatment improves the dom time?
- Q2: Using data in Foo.csv, How much faster is H2 over H1? What percent improvement does H2 offer?
- Q2: Any other interesting things that jump out from the data?

### 0.1.1 Importing Libraries

```
In [4]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
sns.set_style("darkgrid")
import numpy as np
import statsmodels.formula.api as smf
import statsmodels.stats.multicomp as multi
```

### 0.1.2 Functions

```
In [52]: def plot_corr(df,size=10):
    '''Function plots a graphical correlation matrix for each pair of columns in the
    Input:
        df: pandas DataFrame
        size: vertical and horizontal size of the plot
    '''
    corr = df.corr()
    fig, ax = plt.subplots(figsize=(size, size))
    ax.matshow(corr)
    plt.xticks(range(len(corr.columns)), corr.columns);
    plt.yticks(range(len(corr.columns)), corr.columns);
    plt.show()

def BoxPlot(Col1,Col2,Title='',size=10):
    fig, ax = plt.subplots(figsize=(size, size))
    sns.violinplot(x=Col1, y=Col2, showmeans=True,scale="width", inner="quartile")
    plt.title(Title)
```


```
plt.show()
```

### 0.1.3 Importing the CSV files

```
In [96]: BAR=pd.read_csv('data/bar.csv')
        F00=pd.read_csv('data/foo.csv')
        DF_trtmt_dom=BAR[['trtmt','dom']]
        DF_F00=F00[['label','loadTime','fullyLoaded']]

In [54]: BoxPlot(BAR['trtmt'], BAR['dom'],'violet Plot of ....',10)
        DOM_TRM_ON_Mean=DF_trtmt_dom[DF_trtmt_dom['trtmt']=='On']['dom'].mean()
        DOM_TRM_OFF_Mean=DF_trtmt_dom[DF_trtmt_dom['trtmt']=='Off']['dom'].mean()

print ('\nThe average dom time with treatement = {}'.format(DOM_TRM_ON_Mean) )
print ('The average dom time without treatement = {} \n'.format(DOM_TRM_OFF_Mean) )
```



The average dom time with treatement = 19.30002315993797  
The average dom time without treatement = 21.080772215776854

- The difference in the treatment mean does NOT prove the relationship between the treatment and the dom time.
- We now run Anova F-statistic test to identify if the Treatment (categorical predictor variable) associated or related with the dom time (quantitative target variable)?
- To prove the relationship we look for the Anova test F-statistic and Prob (F-statistic) values.

```
In [63]: model = smf.ols(formula='dom ~ trtm', data=BAR)
         results = model.fit()
         FValue=results.fvalue
         FProp=results.f_pvalue
         R_sqr=results.rsquared
         print ('\nF-statistic = {}'.format(FValue))
         print ('Prob (F-statistic) = {}'.format(FProp))
         print ('R-squared = {}'.format(R_sqr))
         #print (results.summary())
```

F-statistic =272.1850607794322  
Prob (F-statistic) =7.063553955518299e-61  
R-squared =0.009053465451638432

- Given that the F-statistic is very large and the Prob (F-statistic) is very small, then, we can say that the change in dom time is related to the treatment but with very low R-squared value = 0.009.

#### 0.1.4 The treatment improves the dom time by 1.78 (21.08 -19.30)

In [ ]:

The t-statistic is a measure of the difference between the two sets expressed in units of standard error. Put differently, it's the size of the difference relative to the variance in the data. A high t value means there's probably a real difference between the two sets; you have "significance". The P-value is a measure of the probability of an observation lying at extreme t-values; so a low p-value also implies "significance." If you're looking for a "statistically significant" result, you want to see a very low p-value and a high t-statistic (well, a high absolute value of the t-statistic more precisely). In the real world, statisticians seem to put more weight on the p-value result.

Let's change things up so both A and B are just random, generated under the same parameters. So there's no "real" difference between the two:

```

In [68]: import numpy as np
         from scipy import stats

         A=DF_trtmt_dom[DF_trtmt_dom['trtmt']=='On']['dom'].tolist()
         B=DF_trtmt_dom[DF_trtmt_dom['trtmt']=='Off']['dom'].tolist()

         stats.ttest_ind(A, B)

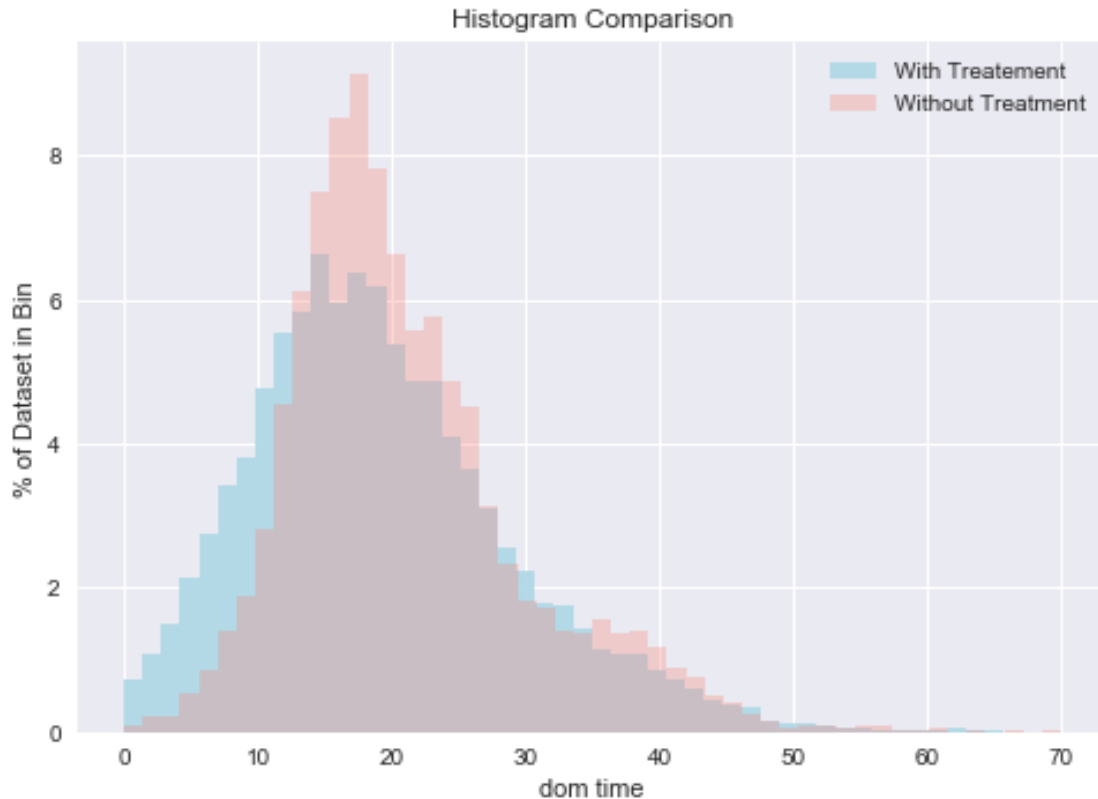
Out[68]: Ttest_indResult(statistic=-16.498032027469961, pvalue=7.063553955755318e-61)

In [81]: import matplotlib.pyplot as plt
         import numpy as np
         np.random.seed(1)

         xweights = 100 * np.ones_like(A) / len(A)
         yweights = 100 * np.ones_like(B) / len(B)

         fig, ax = plt.subplots()
         ax.hist(A, weights=xweights, color='lightblue', alpha=0.9, bins=50)
         ax.hist(B, weights=yweights, color='salmon', alpha=0.3, bins=50)
         ax.legend(['With Treatement', 'Without Treatment'])
         ax.set(title='Histogram Comparison', ylabel='% of Dataset in Bin')
         ax.margins(0.05)
         ax.set_ylim(bottom=0)
         plt.xlabel('dom time')
         plt.show()

```



```
In [86]: F00.columns
```

```
Out[86]: Index(['testId', 'label', 'location', 'completed_time_stamp', 'completed_time',
               'run_no', 'loadTime', 'fullyLoaded', 'bytesIn', 'bytesInDoc',
               'requests', 'requestsDoc', 'TTFB', 'firstPaint', 'render', 'SpeedIndex',
               'docTime', 'ttfb_first_req', 'ssl_first_req', 'dns_first_req'],
              dtype='object')
```

```
In [97]: DF_F00['diff']=DF_F00.apply(lambda row: row[2]-row[1],axis=1)
```

```
/Users/alhaol/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/indexing.html
"""Entry point for launching an IPython kernel.
```

```
In [99]: DF_F00.nunique()
```

```
Out[99]: label          2
         loadTime      742
```

```
fullyLoaded    788
diff            682
dtype: int64
```

```
In [108]: import numpy as np
          from scipy import stats
```

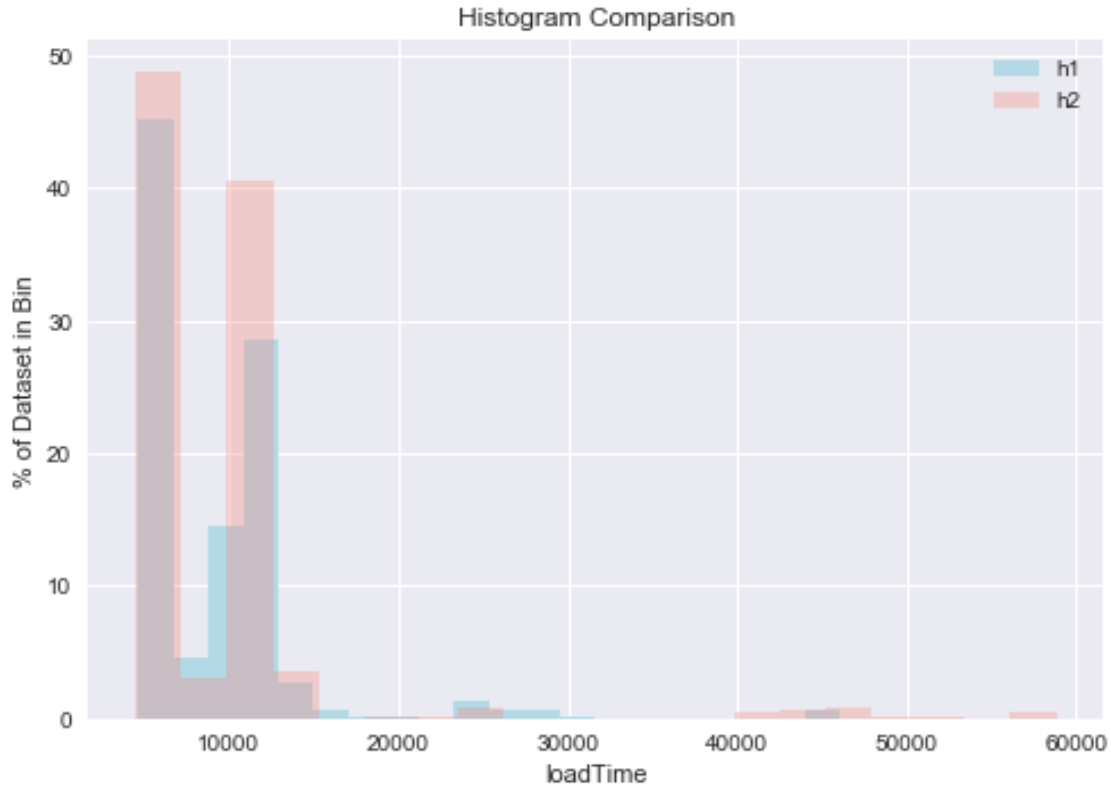
```
A=DF_F00[DF_F00['label']=='h1']['loadTime'].tolist()
B=DF_F00[DF_F00['label']=='h2']['loadTime'].tolist()

stats.ttest_ind(A, B)
```

```
Out[108]: Ttest_indResult(statistic=-0.8736228043200106, pvalue=0.3825426361512243)
```

```
In [110]: xweights = 100 * np.ones_like(A) / len(A)
          yweights = 100 * np.ones_like(B) / len(B)

fig, ax = plt.subplots()
ax.hist(A, weights=xweights, color='lightblue', alpha=0.9, bins=20)
ax.hist(B, weights=yweights, color='salmon', alpha=0.3, bins=20)
ax.legend(['h1', 'h2'])
ax.set(title='Histogram Comparison', ylabel='% of Dataset in Bin')
ax.margins(0.05)
ax.set_ylim(bottom=0)
plt.xlabel('loadTime')
plt.show()
```



```
In [111]: import numpy as np
          from scipy import stats
```

```
A=DF_F00[DF_F00['label']=='h1']['fullyLoaded'].tolist()
B=DF_F00[DF_F00['label']=='h2']['fullyLoaded'].tolist()
```

```
stats.ttest_ind(A, B)
```

```
Out[111]: Ttest_indResult(statistic=-0.81850317851539156, pvalue=0.41327347416106186)
```

```
In [ ]: xweights = 100 * np.ones_like(A) / len(A)
        yweights = 100 * np.ones_like(B) / len(B)
```

```
fig, ax = plt.subplots()
ax.hist(A, weights=xweights, color='lightblue', alpha=0.9, bins=20)
ax.hist(B, weights=yweights, color='salmon', alpha=0.3, bins=20)
ax.legend(['h1', 'h2'])
ax.set(title='Histogram Comparison', ylabel='% of Dataset in Bin')
ax.margins(0.05)
ax.set_ylim(bottom=0)
plt.xlabel('loadTime')
plt.show()
```

**0.2 Unfortunately there is no statistical evidence \* low p-value to prove**

**0.3 that any of them is better, the performance cam from rando**

In [ ]:

In [82]: F00.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 960 entries, 0 to 959
Data columns (total 20 columns):
testId          960 non-null object
label           960 non-null object
location        960 non-null object
completed_time_stamp  960 non-null int64
completed_time   960 non-null object
run_no          960 non-null int64
loadTime        960 non-null int64
fullyLoaded     960 non-null int64
bytesIn         960 non-null int64
bytesInDoc      960 non-null int64
requests        960 non-null int64
requestsDoc     960 non-null int64
TTFB            960 non-null int64
firstPaint      960 non-null int64
render          960 non-null int64
SpeedIndex      960 non-null int64
docTime         960 non-null int64
ttfb_first_req  960 non-null int64
ssl_first_req   960 non-null int64
dns_first_req   960 non-null int64
dtypes: int64(16), object(4)
memory usage: 150.1+ KB
```

In [22]: BAR.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 29794 entries, 0 to 29793
Data columns (total 18 columns):
Unnamed: 0      29794 non-null int64
host            29794 non-null object
page           29794 non-null object
isp            29794 non-null object
browser        29794 non-null object
device         29794 non-null object
trtmt          29794 non-null object
dom            29794 non-null float64
queryString    29794 non-null object
```



```


f1          29794 non-null float64
f2          29794 non-null float64
f3          29794 non-null float64
f4          29794 non-null float64
f5          29794 non-null float64
f6          29794 non-null float64
f7          29794 non-null float64
f8          29794 non-null float64
f9          29794 non-null float64
dtypes: float64(10), int64(1), object(7)
memory usage: 4.1+ MB

```

```

In [75]: plt.matshow(BAR.corr())
         plt.show()


```




```

In [83]:
         plot_corr(BAR,size=10)
         plot_corr(F00,size=10)

```






In [94]: BAR.columns.tolist()

Out[94]: ['host',  
 'page',  
 'isp',  
 'browser',  
 'device',  
 'trtm',  
 'dom',  
 'queryString',  
 'f1',  
 'f2',  
 'f3',  
 'f4',  
 'f5',  
 'f6',  
 'f7',  
 'f8',  
 'f9']

```
In [103]: import seaborn as sns
sns.pairplot(BAR[['dom', 'f1', 'f2', 'f3', 'f4', 'f5', 'f6', 'trtmmt']], hue='trtmmt')
plt.show()
```




```
In [106]: BAR[['dom', 'f1', 'f2', 'f3', 'f4', 'f5', 'f6', 'trtmmt']].head(3)
```

```
Out[106]:
```

	dom	f1	f2	f3	f4	f5	f6	trtmmt
0	43.846811	0.205366	0.817999	0.533166	0.358433	0.069712	0.370256	On
1	34.078843	0.455522	0.793443	0.090781	0.609515	0.746002	0.511265	On
2	10.423054	0.195735	0.491308	0.147366	0.797300	0.409352	0.857075	On

```
In [110]:
```



```
In [113]: #ANOVA F Test
import numpy as np
import pandas as pd
import statsmodels.formula.api as smf
import statsmodels.stats.multicomp as multi
model = smf.ols(formula='dom ~ trtmnt', data=BAR)
results = model.fit()
print (results.summary())
```

#### OLS Regression Results

```
=====
Dep. Variable:          dom    R-squared:                0.009
Model:                  OLS    Adj. R-squared:           0.009
Method:                 Least Squares    F-statistic:          272.2
Date:                   Sat, 09 Dec 2017    Prob (F-statistic):    7.06e-61
Time:                   13:06:43    Log-Likelihood:       -1.0876e+05
No. Observations:       29794    AIC:                  2.175e+05
Df Residuals:           29792    BIC:                  2.175e+05
Df Model:                1
Covariance Type:        nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
--	------	---------	---	------	--------	--------


```
-----
Intercept      21.0808      0.076    277.446      0.000      20.932      21.230
trtmnt[T.On]   -1.7807      0.108    -16.498      0.000      -1.992     -1.569
=====
Omnibus:                3493.413    Durbin-Watson:                1.982
Prob(Omnibus):           0.000    Jarque-Bera (JB):            5196.387
Skew:                    0.873    Prob(JB):                     0.00
Kurtosis:                4.065    Cond. No.                     2.61
=====
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.


```
In [116]: import seaborn as sns
```

```
f, ax = plt.subplots(figsize=(15, 12))
corr = BAR.corr()
sns.heatmap(corr, mask=np.zeros_like(corr, dtype=np.bool), cmap=sns.diverging_palette(
    square=True, ax=ax)
plt.show()
```



```
In [117]: import seaborn as sns
```

```
f, ax = plt.subplots(figsize=(15, 12))
corr = F00.corr()
sns.heatmap(corr, mask=np.zeros_like(corr, dtype=np.bool), cmap=sns.diverging_palette(
    square=True, ax=ax)
plt.show()
```



Thank you for your interest in the Data Scientist position in our Ottawa office. When you have a chance please complete the attached exercise (ideally we would like the exercise completed and returned in 2 days but if you have other obligations please feel free to let me know). If you have any questions please do not hesitate to call or email me at any time. Regards, Nick Russo

Questions for attached Exercise Bar.csv How much faster does the treatment improves the dom time? Foo.csv: How much faster is H2 over H1? What percent improvement does H2 offer? Any other interesting things that jump out from the data.

```
In [65]: DOM_TRM=BAR[['dom','trtmt']]
```

```
In [66]: import numpy as np
```

```
DOM_TRM_ON=DOM_TRM[DOM_TRM['trtmt']=='On']['dom']
DOM_TRM_off=DOM_TRM[DOM_TRM['trtmt']=='Off']['dom']
DOM_TRM_off=pd.sub
```


File "<ipython-input-66-d08cd9db6fd4>", line 4



```
DOM_TRM_ON=DOM_TRM[DOM_TRM['trtmt']=='On']['dom']  
^
```

IndentationError: expected an indented block

```
In [60]: import matplotlib.pyplot as plt  
DOM=DOM_TRM_ON-DOM_TRM_off  
DOM.plot()  
plt.show()
```



```
In [54]: DOM_TRM_off
```

```
Out[54]: 15030
```

```
In [47]: F00[F00['label']=='h1']['fullyLoaded'].mean()
```

```
Out[47]: 9168.816666666668
```

```
In [48]: F00[F00['label']=='h2']['fullyLoaded'].mean()
```

```
Out[48]: 9506.891666666666
```

```
In [49]: Waht to do:
```

- 1) Correlation between Features
- 2) PCA and visulization of the fature
- 3) Random forest predition
- 4) Correlation matrix